Project 1
FYS-STK3155

Magnus Bergkvam

October 10, 2023

1 Abstract

We have in this project fitted various models on data generated from the franke-function, a function which somewhat resembles a landscape, as well as real-world geographical data. In both cases we have used polynomial features up to order 5 related to a position in the "landscape" in some way. For the franke-function we have used both response with and without added noise, and for the landscape data we naturally have some stochastic noise built in. The models we looked at was ridge, lasso and ordinary least squares. For both the mean-squared-error and the R^2 score we saw that for the franke-function without noise the models performed rather similarly, but when we added noise simpler versions of ols and quite heavily penalized ridge/lasso models performed much better than the more complex models, while the same couldn't be said for the data-rich landscape data. Throughout the project we have clearly seen that less complex models often perform better when we have some added noise (at least in the case when we have somewhat limited data), but that this is not the case when there is no noise (or a very high signal to noise-ratio). We have also explored how this result is tied to the bias-variance tradeoff, and manifested the importance this has for choosing a correct model complexity for our given data.

2 Introduction

Linear regression often is seen as a rather simple machine-learning method in itself, but we do have less complex models than this again. Choosing the right model complexity is well known as being an important part in getting the best possible models for prediction, and in a lot of applications linear regression offers just the right amount of complexity needed. Especially with cases where our data is limited, or we have a low signal-to-noise ratio, linear regression often performs very well [1, p. 43].

However, when can it become beneficial to choose even less complex models than this again? This is mainly what this project is about. We will mainly

base this project around the franke-function, which is a 2-dimensional function which somewhat resembles a landscape [2]. On the franke-function we will fit ordinary least squares linear regression with polynomial features up to order 5 (with interactions), and then compare these results (using mainly the mean squared error and R^2 metrics) to ridge and lasso regression on the same data. We will also explore the performances of these models on some real-world geographical data as well. For the franke-function we will be generating two vectors of random points in $[0, 1]$, and use polynomial features with interactions of order 5 for our data, for all our models. We will also try to tie the results we get to the well-known bias-variance tradeoff in machine learning, and on our way there explore various resampling techniques for getting better estimates for model performances.

We now get into a description of the models/methods we have used, before we present the results and finally conclude what we have found out.

3   Methods

3.1   The franke function

The franke-function is a function of two variables with the following definition:

$$f(x, y) = \frac{3}{4} \exp\left(-\frac{(9x-2)^2}{4} - \frac{(9y-2)^2}{4}\right) + \frac{3}{4} \exp\left(-\frac{(9x+1)^2}{49} - \frac{(9y+1)}{10}\right)$$

$$+ \frac{1}{2} \exp\left(-\frac{(9x-7)^2}{4} - \frac{(9y-3)^2}{4}\right) - \frac{1}{5} \exp\left(-(9x-4)^2 - (9y-7)^2\right)$$

We will use this function to generate a response from two explanatory variables x_1 and x_2 both of which will be drawn randomly from $[0, 1]$. To get a visual picture of what the franke-function looks like on $[0, 1] \times [0, 1]$ see figure 1.

Looking at such a plot we can see that the franke-function does somewhat resemble a terrain. We will use the franke-function to generate data, both with and without noise. In the case without noise we will simply get the following response:

$$\mathbf{y} = f(\mathbf{x_1}, \mathbf{x_2})$$

(Here where f is applied element-wise to $\mathbf{x1}$ and $\mathbf{x2}$). With noise we simply get:

$$\mathbf{y} = f(\mathbf{x_1}, \mathbf{x_2}) + \epsilon$$

Where in this case $\epsilon_i \sim N(0, \sigma^2)$. We will use $\sigma = 0.1$ throughout all of our calculations (but this is easily adjustable in the code).

3.2   Geographical data

As mentioned, another source for data which was used for analysis of the different models is some real-world geographical data. The data is stored as a

Figure 1: A surface plot of the franke-function on $[0, 1] \times [0, 1]$

tiff image file and consists simply of a set of numerical values in a square grid, where the numerical value indicates the height of the terrain in some way. Figure 2 contains a plot of the data which we have used for our analysis here. The data is gathered from `https://earthexplorer.usgs.gov/`, and we have here downloaded a region around *Møsvatn Austfjell* in Norway. This data includes a lot of datapoints ($3601 \times 1801 = 6485401$ to be exact), which is quite a lot more than we need to fit a good model. Trying to fit a lot of different models to such a big amount of data will often take a lot of memory and processing power, so for this sake I have throughout our calculations limited the data to look at only the first $400 \times 500 = 20000$ observations, which should be more than enough for our case. When reading the tiff-file we only get a matrix of values out, but we will look at this as an image, and then use the position of the response in the matrix as explanatory variables for our model. In other words, if say the fifth pixel in the fourth row has the value 10^{-5}, the response will obviously be 10^{-5}, while the two explanatory variables will have values 5 and 4. This way we will be able to gather data whose structure closely resembles that of the franke-function. We can then use the same techniques to set up a design-matrix and a response, as we do with the franke-function.

3.3   Data/Data scaling

The data and design matrix used throughout the project is mostly the same for the different models we fit the data to. As previously mentioned, for the franke-function the data we have used is simply a collection of two randomly generated datapoint-vectors $\mathbf{x}_1$ and $\mathbf{x}_2$, each of size 200 (for the most part), and

Figure 2: A surface plot of the landscape data we have used. Keep in mind here that the axis here is simply the index of each datapoint in the tiff-file. This means that the axis here doesn't really tell us much, but shows the overall structure of the terrain nonetheless.

the response is simply the value of the franke-function evaluated at some point (plus some added noise in some cases). When we added noise we have drawn this from the normal distribution with mean 0 and standard-deviance 0.1, which satisfies a common assumption for linear regression models [3, s. 1.4.5]. We have chosen to use a standard-deviance of 0.1 in order to not have too low of a signal-to-noise-ratio (the franke-function usually returns quite low values as one can can see from figure 1). For the geographical data we load the data in from a tiff file of geographical data, but the setup of the design matrix is exactly the same, with the difference being that the x_1s, x_2s and the response have been drawn directly from the tiff-file. For the setup of the design-matrix of for example fifth order polynomial features we have included in the first column the first-order values of the x_1s, the second column the second-order values of x_1s, and so on up to the fifth order. The sixth to tenth column then is the first to fifth order terms of x_2s and then the rest of the $4 \cdot 4 = 16$ columns is for interaction terms between the two variables. The same structure applies of course to other orders as well. One exception to this exact setup of design-matrix is for ordinary least squares where we have added an intercept by setting the first column to consist of only 1s.

Throughout the whole of this project we have scaled the features of the matrix, simply by centering the mean around 0 and scaling by the standard deviance. This way we have that each feature is 1 away from 0 on average. In the case of ordinary least squares regression this scaling is not important for model per-

4

formance, as the model will be able to scale the βs depending on the scaling of the features and give the same minimization of the loss-function (just with different parameters β). However, in the case of ridge/lasso regression which both include a penalty term, it becomes important to scale the data so that the scale of the data doesn't affect how much we penalize the parameters (more about this in the ridge/lasso section at 3.6).

A last thing of note we have used throughout the project is train-test-splitting (except when doing cross-validation). We then split our data into a train set consisting of 80% of the datapoints, and 20% of the data for testing. This lets us evaluate how well our model performs on "new" data, i.e. data the model has not been trained on. The reason this is important is in order for us to recognize overfitted models, which often become the case when we have a higher model complexity [1, s. 7.2]. In such cases we could have a model that predicts the training data very well, but generalizes very poorly to new data. For the franke-function I have written a function *generatedata* which takes hand of all of the mentioned above (I am using Julia for all the code in this project). It is a bit long, but well worth inclusion here as it is central to all of the programs that use the franke-function:

```
1   # Generate a design matrix consisting of random x-s (two explanatory variables)
2   # with a polynomial of a given order (we will use order 5 for our analysis
3   # mainly). Additionally includes options for adding intercept column, noise,
4   # number of observations and random seed.
5   function generatedata(order::Int64; split=true, include_intercept
        =false, add_noise=false, noise_σ=0.1, n=200, custom_seed
        =1000)
6       # Setting the seed for train test splitting and random x1/x2
7       seed!(custom_seed)
8
9       # Generating x-s
10      x1 = rand(n)
11      x2 = rand(n)
12
13      # Creating the design matrix
14      X = generatedesignmatrix(x1, x2, order)
15
16      # Creating the response
17      y = Functions.frankefunction(x1, x2)
18
19
20      # Shuffle before train-test splitting
21      Xs, ys = shufflematrices(X, y)
22
23      # Train-test splitting
24      if !split
25          return standardscale(Xs, Xs), ys
26      end
27      indextosplitat = Int(floor(size(Xs, 1) * 0.8))
28      X_train, X_test = Xs[1:indextosplitat, :], Xs[(indextosplitat
            +1):size(Xs, 1), :]
```

```
29      y_train, y_test = ys[1:indextosplitat, :], ys[(indextosplitat
            +1):size(ys, 1), :]

30

31      # Scaling X_train/X_test
32      # We use the column means from the original X_train to subtract from the
33      # columns in X_test. We must do this before we scale the X_train
34      X_test = standardscale(X_test, X_train)
35      # We then can scale the X_train
36      X_train = standardscale(X_train, X_train)

37

38      if add_noise
39          # Add response with noise
40          y_train += rand(Normal(0, noise_σ), length(y_train))
41          y_test += rand(Normal(0, noise_σ), length(y_test))
42      end

43

44      if include_intercept
45          X_train = [ones(size(X_train, 1)) X_train]
46          X_test = [ones(size(X_test, 1)) X_test]
47      end

48

49      return X_train, X_test, y_train, y_test
50  end
```

Keep in mind here that we are using the features of the training-set to scale the features of the test-set. We have also used some other local functions like **standardscale** (scaling as described before) and **shufflematrices** (randomly shuffles rows in a matrix), which their implementation is rather trivial (but of course feel free to see exactly how they are implemented by looking at the full source-code at: `https://github.com/magnouvean/ml-physics-projects/tree/main/project1/julia`). We also here use **generatedesignmatrix**, which I'll include the implementation of the here, as it is a little more complicated, and very central for all the calculations we have done (including setting up the landscape data):

```
1  # Generate a matrix with features as polynomial terms with interactions
2  function generatedesignmatrix(x1, x2, order)
3      X = zeros((length(x1), 2 * order + (order - 1)^2))
4      for i in 1:order
5          X[:, i] = x1 .^ i
6          X[:, order+i] = x2 .^ i
7      end
8      for i in 1:(order-1)
9          for j in 1:(order-1)
10             X[:, 1+order+i*(order-1)+j] = (x1 .^ i) .* (x2 .^ j)
11         end
12     end
13
14     return X
15 end
```

Shuffling the data is perhaps not nessecary here in the case of the franke-function seeing as the explanatory variables are randomly drawn (but I have done so anyway just out of good practice), however in the case of the landscape data this is very important, as here the data is sorted when reading the data in, and if we then do not shuffle the data before, and we apply train-test-splitting, we will get different regions of our response for the training and testing data, in other words, the model will be trained on data from some area and tested on data from another area, which will then probably give quite poor performance (unless the two regions the train-data and test-data are gathered from are very similar).

3.4 Metrics

We will mainly use two different metrics in our analysis. The perhaps most central one is the mean squared error, and the other one is the R^2 score. Given some response $\mathbf{y}$ with n values and some predictions for $\mathbf{y}$, $\hat{\mathbf{y}}$ the mse is given as: [4]

$$MSE(\mathbf{y}, \hat{\mathbf{y}}) = \frac{1}{n} \sum_{i=1}^{n} (y_i - \hat{y}_i)^2$$

while the R^2 score is given by: [4]

$$R^2(\mathbf{y}, \hat{\mathbf{y}}) = 1 - \frac{\sum_{i=0}^{n-1}(y_i - \hat{y}_i)^2}{\sum_{i=0}^{n-1}(y_i - \overline{y})^2}$$

3.5 Ordinary least squares

The first, and conceptually simplest, model we looked at in this project was the ordinary least squares linear regression. Before we go deeper into this method, denote $\mathbf{X}$ as our design-matrix and $\mathbf{y}$ as our response. We assume the following relationship:

$$\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\epsilon}$$

Where $\boldsymbol{\epsilon}$ is some stochastic noise. We want to include a intercept for our model, and thus we let the first column in the design matrix $\mathbf{X}$ consist of only 1s. Ordinary least squares regression finds the $\boldsymbol{\beta}$ which minimizes the loss-function over the training-data (this optimal value we call $\hat{\boldsymbol{\beta}}$) given like:

$$C(\boldsymbol{\beta}) = (\mathbf{y} - \mathbf{X}_{train}\boldsymbol{\beta})^T(\mathbf{y} - \mathbf{X}_{train}\boldsymbol{\beta})$$

Then we get the least squares estimate for $\hat{\boldsymbol{\beta}}_{ols}$ given by [1, s. 3.2]:

$$\hat{\boldsymbol{\beta}}_{ols} = (\mathbf{X}_{train}^T \mathbf{X}_{train})^{-1} \mathbf{X}_{train}^T \mathbf{y}_{train}$$

We then can make predictions for the training-set by:

$$\hat{\mathbf{y}}_{train} = \mathbf{X}_{train}\hat{\boldsymbol{\beta}}_{ols}$$

and for the test set:

$$\hat{\mathbf{y}}_{test} = \mathbf{X}_{test}\hat{\boldsymbol{\beta}}_{ols}$$

3.6 Ridge and lasso

Ridge and lasso is in lots of ways very similar to ordinary least squares regression. The setup of the model is exactly the same, i.e. we again have:

$$\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\epsilon}$$

The difference arises with the cost function we minimize to estimate the parameters. For ridge we use:

$$C(\boldsymbol{\beta}) = \left(\sum_{i=1}^{n} (y_i - \beta_0 - \beta_1 x_{i1} - \cdots - \beta_p x_{ip}) \right) + \lambda \sum_{i=1}^{p} \beta_i^2$$

and for lasso:

$$C(\boldsymbol{\beta}) = \left(\sum_{i=1}^{n} (y_i - \beta_0 - \beta_1 x_{i1} - \cdots - \beta_p x_{ip}) \right) + \lambda \sum_{i=1}^{p} |\beta_i|$$

for some hyperparameter λ, which determines how much we want to penalize high parameter values. For ridge we get the following estimator [1, s. 3.4]:

$$\hat{\boldsymbol{\beta}}_{ridge} = (X^T X + \lambda I)^{-1} X^T y$$

Also, keep in mind that we do not here penalize the intercept in the loss-function. We often do this because we in general are ok with a somewhat high intercept, because this often lead to better performance. We will therefore when writing our code, fit ridge regression without a intercept (by not including a first column of only 1s), and then based on the other parameters calculate the intercept β_0 which leads to the optimal loss. This we calculate this by the following relationship:

$$\beta_0 = \frac{1}{n} \sum_{i=0}^{n-1} y_i - \frac{1}{n} \sum_{i=0}^{n-1} \sum_{j=1}^{p-1} X_{ij} \beta_j$$

, as proved in [5, s. Further manipulations]. To get our estimate we simply replace our β_is with their estimates. For lasso-regression we will use scikit-learns implementation, which takes care of this for us. Furthermore, from the loss-function expressions we see why it is important to scale the data beforehand. If we for example have a feature that is not scaled and takes very low values, the corresponding β which minimizes the first sum (the ols loss) will typically be quite high. This would then result in this β to be penalized quite a lot. If we then scale the feature to another unit (by for example multiplying with some big number), so that the feature takes much higher values the corresponding β which minimizes the first sum would then become a lot smaller. This would mean that this parameter would be much less penalized. In other words we see that the scaling of the different features will decide how much its corresponding parameter will be penalized, and this is not something we want. In order to hinder this happening it is important to scale the features somewhat alike, so that the scale the features does not determine which parameters will be shrunk a lot, and which will be shrunk less.

3.7 The bias variance tradeoff

Central in this project will be the bias variance tradeoff, as this is the reason why ridge and lasso often can give better predictions than ordinary least squares. We will here derive the bias-variance equation. Consider the model $\mathbf{y} = f(\mathbf{x}) + \boldsymbol{\epsilon}$ where $E(\boldsymbol{\epsilon}) = 0$ and $Var(\boldsymbol{\epsilon}) = \sigma^2 \mathbf{I}$, for some function f. The job in machine-learning is to estimate this f. We call $\hat{f}$ some estimate of this function f. We then get $\hat{\mathbf{y}} = \hat{f}(\mathbf{x})$, which is some prediction. We consider the squared error of one such prediction $(y - \hat{y})^2$. We now want to find an expression for this mean-squared error, however this is a little difficult to calculate as both terms here are stochastic. We therefore calculate the expectance of this expression, in which we get:

$$
\begin{aligned}
E((y - \hat{y})^2) &= E((f(x) + \epsilon - \hat{f}(x))^2) \\
&= E((\epsilon + (f(x) - \hat{f}(x)))^2) \\
&= E(\epsilon^2 + 2 \cdot \epsilon \cdot (f(x) - \hat{f}(x)) + (f(x) - \hat{f}(x))^2) \\
&= E(\epsilon^2) + 2 \cdot E(\epsilon) \cdot E(f(x) - \hat{f}(x)) + E((f(x) - \hat{f}(x))^2) \\
&= E((\epsilon - E(\epsilon))^2) + E((f(x) - \hat{f}(x))^2) \\
&= \sigma^2 + E((f(x) - \hat{f}(x))^2)
\end{aligned}
$$

Here we have used independence of ϵ from $f(x) - \hat{f}(x)$ which stems from the fact that we may assume that $\hat{f}(x)$ and ϵ are independent. We now focus on calculating the second part of this, i.e. $E((f(x) - \hat{f}(x))^2)$. We have:

$$
\begin{aligned}
&E\left[(f(x) - \hat{f}(x))^2\right] \\
&= E\left[((f(x) - E(\hat{f}(x))) + (E(\hat{f}(x)) - \hat{f}(x)))^2\right] \\
&= E\left[(f(x) - E(\hat{f}(x)))^2 + 2 \cdot (f(x) - E(\hat{f}(x))) \cdot (E(\hat{f}(x)) - \hat{f}(x)) + (E(\hat{f}(x)) - \hat{f}(x))^2\right] \\
&= E\left[(f(x) - E(\hat{f}(x)))^2\right] + 2 \cdot E\left[(f(x) - E(\hat{f}(\mathbf{x}))) \cdot (E(\hat{f}(x)) - \hat{f}(x))\right] + E\left[(E(\hat{f}(x)) - \hat{f}(x))^2\right] \\
&= (f(\mathbf{x}) - E(\hat{f}(x)))^2 + 2 \cdot \left(f(x) - E(\hat{f}(x))\right) \cdot E\left[(E(\hat{f}(x)) - \hat{f}(x))\right] + E\left[(E(\hat{f}(x)) - \hat{f}(x))^2\right] \\
&= (f(\mathbf{x}) - E(\hat{f}(x)))^2 + 2 \cdot \left((f(x) - E(\hat{f}(x)))\right) \cdot (E(\hat{f}(x)) - E(\hat{f}(x))) + E\left[(E(\hat{f}(x)) - \hat{f}(x))^2\right] \\
&= (E(\hat{f}(x)) - f(x))^2 + E\left[(\hat{f}(x) - E(\hat{f}(x)))^2\right]
\end{aligned}
$$

Typically in machine learning we call the first term here the squared bias of the model, and the second term the variance. In total we get:

$$
E((y - \hat{y})^2) = \sigma^2 + (E(\hat{f}(x)) - f(x))^2 + E\left[(\hat{f}(x) - E(\hat{f}(x)))^2\right] = \sigma^2 + Bias(\hat{f}(x))^2 + Var(\hat{f}(x))
$$

In vector form we then see that:

$$E((\mathbf{y} - \hat{\mathbf{y}})^T(\mathbf{y} - \hat{\mathbf{y}})) = \sigma^2 + Bias(\hat{f}(\mathbf{x}))^2 + Var(\hat{f}(\mathbf{x}))$$

We see that the expected squared prediction error is given as a sum of the variance of the noise σ^2, the bias of our model, and the variance of our model. Here σ^2 is an irreducible part we cannot control by the choice of our model, while the bias and variance is dependent on the type of model that we use [1, s. 2.9]. We can see here that the bias is the difference between our models' expectation and the actual expectation $f(\mathbf{x})$. The variance says how much the model varies in predictions when trained on different data. Typically we have that the more complex the model we have the lower the bias will be, but we also typically have a higher variance [1, fig. 2.11]. We also have that the more observations we have the lower the variance usually becomes. Both of these concepts are very central for all our results. We now go more into detail of this for the models we use in the project.

3.8   Bias and variance of ols/ridge/lasso

To explain the bias and variance of the different models, we will dedicate this section to calculate bias and variance for the parameters in ols and ridge regression. Lasso regression is more complicated to do these analysis for as we do not have an analytical expression for $\hat{\boldsymbol{\beta}}$, so we will skip this here. For all our models we usually assume that:

$$y_i = \mathbf{X}_{i*}\boldsymbol{\beta} + \epsilon_i$$

With $\epsilon_i \sim N(0, \sigma^2)$, and $\mathbf{X}_{i*}$ being the i-th row in the design-matrix. We will start by looking at some properties of our response $\mathbf{y}$. Looking first at the expectance of y_i it is easy then to see:

$$
\begin{aligned}
E(y_i) &= E\left[\mathbf{X}_{i*}\boldsymbol{\beta} + \epsilon_i\right] \\
&= E\left[\mathbf{X}_{i*}\boldsymbol{\beta}\right] + E\left[\epsilon_i\right] \\
&= \mathbf{X}_{i*}\boldsymbol{\beta} + 0 \\
&= \mathbf{X}_{i*}\boldsymbol{\beta}
\end{aligned}
$$

(Here we have used the fact that $\mathbf{X}_{i*}\boldsymbol{\beta}$ is non-stochastic and $E(\epsilon_i) = 0$). Calculating the variance of y_i we have:

$$
\begin{aligned}
Var(y_i) &= Var(\mathbf{X}_{i,*}\boldsymbol{\beta} + \epsilon_i) \\
&= Var(\epsilon_i) \\
&= \sigma^2
\end{aligned}
$$

(Again here we have just used that $\mathbf{X}_{i*}\beta$ is non-stochastic). This is enough to conclude that $y_i \sim N(\mathbf{X}_{i*}\beta, \sigma^2)$. To connect this to the bias-variance tradeoff we analyze the expectance and variance of our $\hat{\boldsymbol{\beta}}$ for at least ols and ridge. We start by looking at the ols case. In order to elegantly calculate this we use

two somewhat basic results, which I have proved in the appendix. Let A be a non-stochastic $(p \times n)$ matrix and $\mathbf{Y}$ be a stochastic $(n \times 1)$ vector. We then have the following two results:

$$E(A\mathbf{Y}) = AE(\mathbf{Y})$$

and

$$Var(A\mathbf{Y}) = AVar(\mathbf{Y})A^T$$

Using these results and using that for ols we have $\hat{\beta} = (\mathbf{X}^T\mathbf{X})^{-1}\mathbf{X}^T\mathbf{Y}$ we get:

$$
\begin{aligned}
E(\hat{\boldsymbol{\beta}}) &= E((\mathbf{X}^T\mathbf{X})^{-1}\mathbf{X}^T\mathbf{Y}) \\
&= (\mathbf{X}^T\mathbf{X})^{-1}\mathbf{X}^T E(\mathbf{Y}) \\
&= (\mathbf{X}^T\mathbf{X})^{-1}\mathbf{X}^T\mathbf{X}\boldsymbol{\beta} \\
&= \boldsymbol{\beta}
\end{aligned}
$$

and

$$
\begin{aligned}
Var(\hat{\boldsymbol{\beta}}) &= Var((\mathbf{X}^T\mathbf{X})^{-1}\mathbf{X}^T\mathbf{Y}) \\
&= (\mathbf{X}^T\mathbf{X})^{-1}\mathbf{X}^T Var(\mathbf{Y}) \left((\mathbf{X}^T\mathbf{X})^{-1}\mathbf{X}^T\right)^T \\
&= (\mathbf{X}^T\mathbf{X})^{-1}\mathbf{X}^T \sigma^2 I \mathbf{X} \left((\mathbf{X}^T\mathbf{X})^{-1}\right)^T \\
&= \sigma^2 (\mathbf{X}^T\mathbf{X})^{-1}\mathbf{X}^T\mathbf{X} \left((\mathbf{X}^T\mathbf{X})^T\right)^{-1} \\
&= \sigma^2 (\mathbf{X}^T\mathbf{X})^{-1}\mathbf{X}^T\mathbf{X}(\mathbf{X}^T\mathbf{X})^{-1} \\
&= \sigma^2 (\mathbf{X}^T\mathbf{X})^{-1}
\end{aligned}
$$

For ridge we have that $\hat{\boldsymbol{\beta}}_{ridge} = (\mathbf{X}^T\mathbf{X} + \lambda\mathbf{I}_{pp})^{-1}\mathbf{X}^T\mathbf{Y}$. We now calculate the expectance of this:

$$
\begin{aligned}
E(\hat{\boldsymbol{\beta}}_{ridge}) &= E((\mathbf{X}^T\mathbf{X} + \lambda\mathbf{I}_{pp})^{-1}\mathbf{X}^T\mathbf{Y}) \\
&= (\mathbf{X}^T\mathbf{X} + \lambda\mathbf{I}_{pp})^{-1}\mathbf{X}^T E(\mathbf{Y}) \\
&= (\mathbf{X}^T\mathbf{X} + \lambda\mathbf{I}_{pp})^{-1}\mathbf{X}^T\mathbf{X}\boldsymbol{\beta}
\end{aligned}
$$

and for the variance:

$$
\begin{aligned}
Var(\hat{\boldsymbol{\beta}}_{ridge}) &= Var((\mathbf{X}^T\mathbf{X} + \lambda\mathbf{I}_{pp})^{-1}\mathbf{X}^T\mathbf{Y}) \\
&= (\mathbf{X}^T\mathbf{X} + \lambda\mathbf{I}_{pp})^{-1}\mathbf{X}^T Var(\mathbf{Y})((\mathbf{X}^T\mathbf{X} + \lambda\mathbf{I}_{pp})^{-1}\mathbf{X}^T)^T \\
&= (\mathbf{X}^T\mathbf{X} + \lambda\mathbf{I}_{pp})^{-1}\mathbf{X}^T \sigma^2 \mathbf{I}_{pp}\mathbf{X}((\mathbf{X}^T\mathbf{X} + \lambda\mathbf{I}_{pp})^{-1})^T \\
&= \sigma^2 (\mathbf{X}^T\mathbf{X} + \lambda\mathbf{I}_{pp})^{-1}\mathbf{X}^T\mathbf{X}((\mathbf{X}^T\mathbf{X} + \lambda\mathbf{I}_{pp})^{-1})^T
\end{aligned}
$$

Using these results it is easy to see that under the assumption that $\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\epsilon}$ which we have assumed, that the ordinary least squares model is unbiased

(because $Bias(\hat{f}(\mathbf{x})) = (E(\hat{f}(\mathbf{x})) - f(\mathbf{x})) = \mathbf{x}\boldsymbol{\beta} - \mathbf{x}\boldsymbol{\beta} = 0$), while this is not the case for ridge-regression. However we see that the variance of the ridge-model has potential for lower variance than the ols if we set λ big enough. This does allow for ridge to outperform ols, since the bias-variance tradeoff states that the expected prediction error is the sum of the bias and variance. Much of this effect also applies for lasso regression as well as this also has a higher bias and lower variance than ols.

3.9 Bootstrap

Since this project essentially boils down to comparing metrics for different models on some data it becomes important that these metrics which we calculate are reliable. In the case of the franke-function we have used 200 datapoints, which we have split into a training-set and test-set. We have then trained the data on the training-set and then evaluated the trained model on the test-set, however since the test-set is only 20% of our datapoints, we have that this test-set is rather small. This can make the test-error which we achieve quite specific to the model which we have trained, and our metrics may not be that reliable. In order to combat this we will use resampling techniques. One way which do this by is using bootstrap. We will use bootstrapping to estimate the MSE and R^2 scores on the test-data. Thus we will first have to split the data into a training and testing-set like before. Then we draw bootstrap samples (i.e. we draw n observations/rows from the $n \times p$ design-matrix X_{train} with replacement) from the training-data and fit the model on this data and evaluate the metrics on the test-data. We do this B times and take the mean of these metrics to get our final estimates. In this way we manage to get the MSE and R^2 scores not based on only one model fit, but from B model fits from B training datasets. This often will give us a better estimate for the metrics because it does not require assumptions on error terms, which causes bootstrap to give us more reliable results when we have a smaller sample size (which we do in this case, especially for the franke-function) [6, s. 5.3]. One perhaps weakness with our approach here is that our estimates could then become dependent on the test-data we have picked. This is just something we will have to live with.

3.10 Cross-validation

K-fold cross-validation is another resampling-technique which also can better our estimates for the various metrics. Unlike with bootstrapping we don't need to split the data into a training-set or test-set anymore. K-fold cross-validation works by splitting the data into K folds of as equal sizes as possible. Now denote $\mathbf{X}^{-i}$ as the design-matrix with fold i excluded for some $1 \leq i \leq K$, and $\mathbf{X}^i$ as the data for fold i of our design-matrix. For cross-validation we do for $i = 1, \ldots, K$:

1. Fit our model on $\mathbf{X}^{-i}$

2. Evaluate our metric on $\mathbf{X}^i$

We then, much like bootstrapping average over all the metrics to get our final metric. The reason this improves the estimate for the metrics is that since every fold is used for evaluating the metric, that every observation also is used in the estimate for our final metric. This is in contrast to just simple train-test-splitting where the metric is only evaluated using the rather small test-set. This should give us a better estimate for the true MSE and R^2.

4 Results

4.1 Franke-function

4.1.1 Ordinary least squares

We have located the code for performing ordinary least squares regression on the franke-function with polynomial features up to order 5 in the **LinearRegression.jl** on the project repository [7]. Running this we can extract the following optimal metrics:

	Without noise		With noise	
	MSE	R̂2	MSE	R̂2
TRAIN	0.001046	0.985926	0.008400	0.905406
TEST	0.005326	0.935900	0.018911	0.804954

Where the optimal MSE/R^2 metrics were acquired for polynomial order 3 (both with and without noise). We see there is quite good performance on the training-set and test-set when noise is not included. When we include noise-terms the model-performance on the training data is still quite good, but the performance on the test-data does take a hit (so do the training metrics as well, but not as much). If we look at higher order polynomials from figure 3, where we have no added noise, we can see that the higher orders (in particular order 5) performs quite comparatively to the minimum of order 3 (for new data of course). However in the case where we have added noise, which we have in figure 4, this is not the case, as here the fifth order polynomial model performs very poorly, much worse than order 3 in fact. The fact that this happens only when we have added noise is a sign that we more quickly overfit the models when we have more stochastic noise. Additionally if we look at the average size of values our parameters β takes, as shown in figure 5, we see that the parameters take bigger and bigger values as we increase the polynomial order. Here we also see that the average β values in the case where we have noise is considerably bigger than the case where we do not have noise. Typically the variance increases as the β values get bigger as changes to the data will then typically change the parameters by a lot when the β values are big, especially in the case where we have little data. All this indicates that the model variance is bigger for higher order polynomials, and from what we know from the bias-variance-tradeoff, this is probably the explanation why we get so poor results on predictions on new data when we increase the polynomial order. This is also probably the reason why the performances in the higher order polynomial cases is so much worse when we have added noise.

Figure 3: A plot of the order of polynomial features against the MSE and R^2 for the ols model fit. Here are using a response without stochastic noise.


Figure 4: A plot of the order of polynomial features against the MSE and R^2 for the ols model fit. Here are using a response with stochastic noise.

Figure 5: The average absolute value of our estimated βs for each polynomial order, in a ols model.

4.1.2 Ridge

If we run the **Ridge.jl** program in the project repository [7] we fit the data with fifth order polynomial features (we skip orders lower here) to ridge models with different penalization parameter λ. From that we can extract the following results:

| | Without noise | | With noise | |
	MSE	R^2	MSE	R^2
TRAIN	0.001046	0.985926	0.008400	0.905406
TEST	0.003592	0.956765	0.017352	0.821028

With the optimal λ for test-data without noise being $1.659587 \cdot 10^{-2}$ and when we include noise it becomes 1.513561. To see more in detail how the MSE and R^2-score varies with different value of the hyperparameter λ, see figure 9 and 10. We see that for the training-data that ridge and ols seemingly performs the same. This is because the best MSE here comes from the lowest value of λ, which is so low that they pretty much give the same fit. Here we see that we actually get better performance on new data than the ols case both in the case where we have not added noise, and when we have. In both cases these differences are quite marginal compared to the optimal valued ols case, but still better. Keep in mind here that we have used fifth order polynomial features, so if we compare it to the ols case with fifth order polynomial features we get a lot better results, especially in the case where we add noise. If we look at figure 8 we see that the average absolute size of our parameters shrink as we increase λ, and hence the variance of the model also probably will. The fact that we manage to reduce the variance in this way is likely the reason why our ridge model with
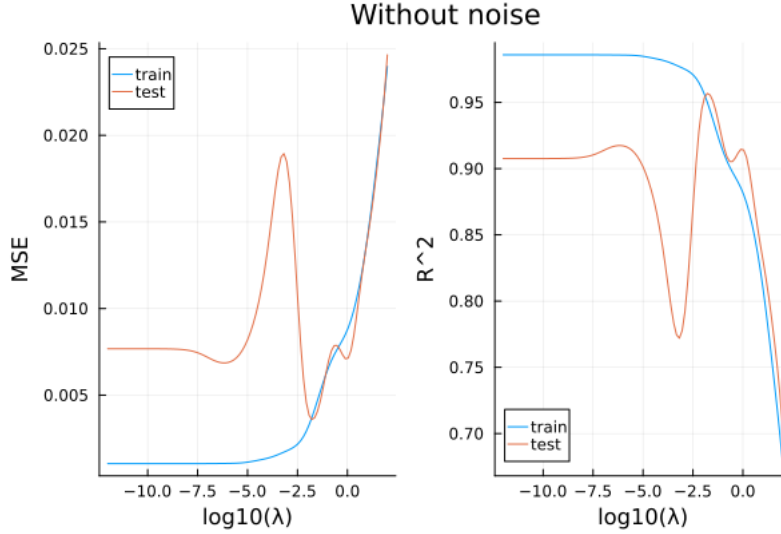
15

Figure 6: A plot of the λ in a ridge-model without noise on log10-scale plotted against the MSE and R^2-score. We see that performances of these metrics goes a bit up and down for the different values of λ.

fifth order polynomial features vastly outperforms the ols model with the same features, at least on the test-data when we have some added noise.

4.1.3 Lasso

The **Lasso.jl** on the project repo [7] runs lasso regression on the franke-function data and calculates the MSE and R^2. Running this we can extract:

	Without noise		With noise	
	MSE	R$\hat{2}$	MSE	R$\hat{2}$
TRAIN	0.003280	0.955866	0.011319	0.872539
TEST	0.002737	0.967054	0.016451	0.830329

With the optimal λ for the training-data being $1.3804 \cdot 10^{-5}$, and for the test data being $9.1201 \cdot 10^{-4}$.

In this case we get an even better performance overall than using ridge in the case of new data, and therefore also better than the ols case, both with and without noise in the response. One thing we notice here with the figures (9 and 10) is that the MSE and R^2 scores completely flatten out after a while. This is because when we increase the λ enough all the parameters will be 0, and then all the metrics will stay the same from hereon naturally. We also see that unlike the ridge-case the MSE and R^2 swings much less in general, and almost looks to be increasing the whole way, with only a small local/global minimums.

Figure 7: A plot of the λ in a ridge-model with λs on log10-scale plotted against the MSE and R^2. We see a somewhat similar relationship as for the case with noise, but with greater scales on the y-axis for MSE.


Figure 8: The λ plotted against the average absolute size of our parameter β. We expectedly get a shrinking relationship as we increase λ.
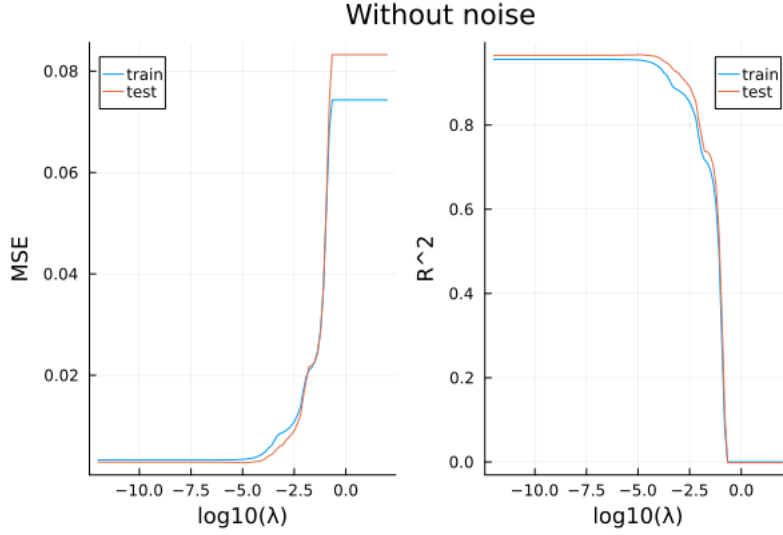
Figure 9: A plot of the λ in a lasso-model with λs on log10-scale plotted against the MSE and R^2. Unlike the Ridge-case this completely flattens out at the end. This is due to the fact that for such high values of λ the model-parameters which minimize the loss is simply with all the βs being 0.


Figure 10: A plot of the λ in a lasso-model with λs on log10-scale plotted against the MSE and R^2. We see the same flattening as in figure 9.

4.1.4   Performance with more complex models

Up until now we have used only polynomial features of order lower than 5. The **Resampling.jl** and **ResamplingCV.jl** on the project repo [7] explores fitting the franke-function to data with more polynomial degrees, with the metrics calculated using bootstrap and cross-validation respectively. Running these programs we gather some plots of the mean squared errors plotted against the model complexity, this time in the form of orders of polynomial features. Figure 11 shows the mse using ols, calculated using bootstrapping, while figure 12, 13 and 14 show the same graphs for ols, ridge and lasso respectively, these calculated using cross-validation. Keep in mind that here we have just set some fixed value for λ for both ridge and lasso, based on what gave good results on fifth order polynomial features in previous calculations. Looking only at the numerical results from the cross-validation for ols on the test-data we got:

Order	2	3	4	5	6	7	8	9	10	11	12	13	14
MSE	0.1822	0.1746	0.1756	0.1796	0.1896	0.2394	0.4217	0.8437	1.1845	1.6786	2.0402	3.2349	4.3811

Similarly for ridge we got:

Order	2	3	4	5	6	7	8	9	10	11	12	13	14
MSE	0.0172	0.0080	0.0058	0.0054	0.0055	0.0099	0.0147	0.02334	0.0364	0.0568	0.0876	0.1290	0.1830

And for lasso:

Order	2	3	4	5	6	7	8	9	10	11	12	13	14
MSE	0.0169	0.0076	0.0061	0.0044	0.0036	0.0033	0.0032	0.0030	0.0029	0.0027	0.0026	0.0025	0.0025

Here we can see the bias-variance tradeoff in full effect. When we include higher order polynomial features, we also get a higher model-complexity. This means our models will have lower bias, but also much higher variance. The expected prediction error being the sum of these two typically means that our model-performances follows a sort of U-shape plotted against model-complexity (this we have already seen for the ols and ridge MSE/R^2 plots in particular), where the optimal value has a somewhat low bias and a somewhat low variance, but also typically not at the minimum of either. This is what we see in general throughout our results here (see figure 11 and 13 in particular). Here the exception seems to be the Lasso, which seems to only get a lower and lower MSE as we increase the polynomial order. This is probably partially because lasso is able to perform its own kind of model-selection, where the parameter corresponding to some of the parameters will be 0. However if we add more and more polynomial features and still keep the λ fixed, we would eventually get an increase in the MSE as the lasso would not be able to reduce the variance enough to keep up with the more and more complex features.

4.2   Landscape data

Running the **Landscape.jl** on the project repo [7] we perform cross-validation on the various models to estimate MSE, and choose the optimal hyperparam-

Figure 11: The MSE calculated using bootstrap plotted against the increasing orders of polynomial features. For the training we unsurprisingly get a lower and lower MSE as we increase the orders of polynomials, but for the testing data we get at first a lower MSE by increasing the polynomials up until about order 6 where it after this increases. This is to be expected from what we know of the bias-variance tradeoff.


Figure 12: Side by side cross-validated estimates of the mean squared error of ordinary least squares models with varying orders of features with $k = 5$ and $k = 10$ folds. We see that the mean squared error for the test data rapidly increases after polynomial order ≈ 6. Since both of the results for $k = 5$ and $k = 10$ seem to be very similar we can be confident that we do not need to increase the amount of folds here to get a more reliable result - the result has already converged.

Figure 13: Side by side cross-validated estimates of the mse, this time for ridge models with varying degrees of polynomial features and $k = 5$ and $k = 10$ this time as well. We get again very similar results for the two sizes of folds. Here we see clearer that the test MSE decreases at first as we get higher order polynomial features, before it increases again as we get more and more complex models. We also notice that the scale here is much smaller than in the ols case, likely due to ridge being a less complex model.


Figure 14: Side by side cross-validated estimates of the mse, finally this time for lasso models with varying degrees of polynomial features and $k = 5$ and $k = 10$ folds this time as well, with the same result for the two different fold sizes. Interesting here is that we actually get better and better performance as we increase the polynomial order it could look like. Keep in mind though that we only have tested up to order 14. If we were to increase this even more, we likely would get problems so long as the λ parameter stays the same.

eters. We here kept to polynomial features of order 5 which seemed to get us quite good performance across the models for the franke-function (perhaps with the exception of ols, but keep in mind that we have much more data here). When using the ols and calculating the MSE using cross-validation we get a MSE value of $\approx 1.6075 \cdot 10^{-5}$, and when evaluating on the separate test-data we get $\approx 1.5878 \cdot 10^{-5}$. For ridge regression we use the cross-validation to determine the optimal λ we get this to be $\approx 4.6416 \cdot 10^{-8}$, with a corresponding MSE to of $\approx 1.6 \cdot 10^{-5}$, much like the ols case. Evaluated on the test-set for this model we get a MSE of $\approx 1.5878 \cdot 10^{-5}$. For lasso we get the optimal λ to be 10^{-10}, which is the smallest possible value, with a corresponding MSE to be $\approx 2.4 \cdot 10^{-5}$. On the test-data this MSE was $\approx 2.37 \cdot 10^{-5}$.

What is noticeable here is that we get the best λ values for prediction for both the ridge and lasso to be very low, meaning we do not get a lot of reduction in the parameters, and therefore also not much reduction in model variance. This is likely because we have a lot more data-points in this case so the model-variances become rather small in all cases, so there is not much to gain from reducing this further, and therefore the models with lower bias will give the best predictions (which is the ols).

5   Analysis

As we can see from the results, in the case of when we have data without noise (as in the franke-function with no noise), we have quite comparably good results when fitting the different models to the data, with lasso performing just about better than ridge, while ols performing the worst, but still quite good. When we add noise to the data all the models perform worse, but lasso and ridge manage to still come up with quite good performance (lasso still performed the best), while higher order polynomial featured ols performed much much worse than before, but lower order ols models performed comparable to ridge/lasso still. What ridge and lasso models with quite big values of λ and lower order polynomial ols have in common is that they are all comparatively not that complex models. The reason these models performed so good is likely because, even though the model bias is somewhat bigger, the variance will be much smaller, and from the bias-variance-tradeoff we know that it is the sum of these which needs to be low for the model to make good predictions. We have seen this also when trying out different polynomial orders and seeing the MSE values these give. In general at least for ridge and ols we see that both calculating the MSE using bootstrap and cross-validation that we get slight decreases in MSE when increasing the polynomial orders, before the MSE eventually increase again with the polynomial order as the models get more complex. Looking at the overall performances for the different models on the different polynomial orders we see that they all work the best on quite different orders. For the ridge one would get the best performance around 5-th order polynomial features, while for ols this is closer to 3-rd order, while for lasso one can go a lot higher than both of the ols and ridge. For the landscape data we saw that we got a similar performance

for the ols and ridge, while lasso all of a sudden performed the worst. This difference is likely related to a different signal-to-noise-ratio than before, but perhaps even more important, due to the fact that we then had many more observations. With more observations we do not overfit like we did before, and the variance of the model becomes lower than when we have less data. This makes linear regression a more viable choice seeing as this often has considerably lower bias.

At last, one thing to take note of is that throughout some of the code (in particular the parts that have to do with ridge/lasso using the franke-function), we have fitted many models with different hyperparameters and then chosen the model which gave the best result for our MSE metric. This MSE metric however would be too optimistic for the MSE on the true data [1, s. 7.2]. Here we could further improve the estimates by generating a separate test-set, which we evaluate our selected best model on. Keep in mind though that we did do this for the landscape data, so our metrics here should be rather fine. Another place we could improve is to test performances for more penalization terms for ridge and lasso. We used a limited set here in order to speed up the calculations, and one could perhaps test a few more in order to see if model-performances becomes better.

6 Conclusion

We have in this project seen how ols, ridge and lasso performs on data generated from the franke-function and landscape-data. We see that for data without noise and quite low amount of observations that all these three different models performed somewhat similarly, but when we added noise that the less complex models, i.e. ridge and lasso with somewhat sizeable λ parameter, and ols with lower order polynomial features performed better. This clearly is related to the bias-variance-tradeoff, which we have derived earlier. However in the case of the landscape-data, where we had a lot more observations and polynomial features of order five, the ols performed much better, and the ridge/lasso with λ parameters of any meaning performed worse. Over the two datasets we have manifested the importance of choosing models of right model-complexity in order to get the best predictions. We also saw that using models with lower variance became important when we had few observations and a sizeable amount of noise, but less important when we have a good amount of observations and no/tiny stochastic noise in the data.

7 Appendix

7.1 Project code (github)

All the Julia programs described in the report and with the full source code can be found at: `https://github.com/magnouvean/ml-physics-projects/tree/main/project1/julia`

7.2 Expectance and variance with matrices

Let A be a non-stochastic $p \times n$ matrix and $\mathbf{X}$ be a stochastic $n \times 1$ vector. We have:

$$E(A\mathbf{X}) = (E(A\mathbf{X})_1, \ldots, E((A\mathbf{X})_n))^T$$

We then have:

$$E(A\mathbf{X})_i = E((A\mathbf{X})_i) = E(\sum_{k=1}^{n} A_{ik}\mathbf{X}_k) = \sum_{k=1}^{n} A_{ik}E(\mathbf{X}_k) = A_{i*}E(\mathbf{X})$$

Hence we see that $E(A\mathbf{X}) = AE(\mathbf{X})$. Note also that from this it follows that $E(A\mathbf{X})^T = (AE(\mathbf{X}))^T = E(\mathbf{X})^T A^T$ (we will use this for calculating the variance expression).

We now want to look at the case of the variance. We first calculate an intermediate expression. Let A as before be a non-stochastic $p \times n$ matrix and $\mathbf{B}$ a stochastic $n \times n$ matrix. We want to find:

$$E(A\mathbf{B}A^T)$$

We have:

$$\begin{aligned}
E(A\mathbf{B}A^T)_{ij} &= E((A\mathbf{B}A^T)_{ij}) \\
&= E(\sum_{k=1}^{n} A_{ik}(\mathbf{B}A^T)_{kj}) \\
&= \sum_{k=1}^{n} A_{ik}E((\mathbf{B}A^T)_{kj}) \\
&= \sum_{k=1}^{n} A_{ik}E(\sum_{l=1}^{n} \mathbf{B}_{kl}A_{lj}^T) \\
&= \sum_{k=1}^{n} A_{ik} \sum_{l=1}^{n} A_{lj}^T E(\mathbf{B}_{kl}) \\
&= (AE(\mathbf{B})A^T)_{ij}
\end{aligned}$$

Hence we see that $E(A\mathbf{B}A^T) = AE(\mathbf{B})A^T$. We are now ready to find the variance:

$$\begin{aligned}
Var(A\mathbf{X}) &= E((A\mathbf{X} - E(A\mathbf{X}))(A\mathbf{X} - E(A\mathbf{X}))^T) \\
&= E((A\mathbf{X})(A\mathbf{X})^T - A\mathbf{X}E(A\mathbf{X})^T - E(A\mathbf{X})(A\mathbf{X})^T + E(A\mathbf{X})E(A\mathbf{X})^T) \\
&= E(A\mathbf{X}\mathbf{X}^T A^T - A\mathbf{X}E(\mathbf{X})^T A^T - AE(\mathbf{X})\mathbf{X}^T A^T + AE(\mathbf{X})E(\mathbf{X})^T A^T) \\
&= E(A(\mathbf{X}\mathbf{X}^T - \mathbf{X}E(\mathbf{X})^T - E(\mathbf{X})\mathbf{X}^T + E(\mathbf{X})E(\mathbf{X})^T)A^T) \\
&= AE((\mathbf{X} - E(\mathbf{X}))(\mathbf{X} - E(\mathbf{X}))^T)A^T \\
&= AVar(\mathbf{X})A^T
\end{aligned}$$

Where we have used the previous result for the last equality.

References

[1] Trevor Hastie, Robert Tibshirani, Jerome H Friedman, and Jerome H Friedman. *The elements of statistical learning: data mining, inference, and prediction*, volume 2. Springer, 2009.

[2] Franke's function. `https://www.sfu.ca/~ssurjano/franke2d.html`, 2023. [Online; accessed 28-September-2023].

[3] Kevin P Murphy. *Machine learning: a probabilistic perspective*. MIT press, 2012.

[4] Week 34. `https://compphysics.github.io/MachineLearning/doc/LectureNotes/_build/html/week34.html`, 2023.

[5] Week 35. `https://compphysics.github.io/MachineLearning/doc/LectureNotes/_build/html/week35.html`, 2023.

[6] Resampling methods. `https://compphysics.github.io/MachineLearning/doc/LectureNotes/_build/html/chapter3.html`, 2023.

[7] Project 1. `https://github.com/magnouvean/ml-physics-projects/tree/main/project1`, 2023.